

Code performance analysis and improvement by optimising compilation : Case study

Pierre Tassel under the supervision of Pr. Touati

Département of Computer Science
University of Nice Sophia-Antipolis

Student in Master of Computer Science

July 4, 2018

Outline

- 1 The program
- 2 Experimental approach
- 3 Analysis and optimization of sequential performances
- 4 Adding parallelism
- 5 Simple algorithmic modification
- 6 Conclusion
- 7 Perspectives

The program

MinMax for the Awale game

- AI for Awale's game with 6 row and 4 seed per row.
- Written in C++ by Pr Régis.



Figure: The Awale's game.

MinMax with Alpha Beta pruning




Experimental approach

Hardware and software environment

- *Processeur Scaling* desactivated, Intel core i7 8 core 2,5GHz processor, 8Go DDR3, 256Go of SSD.
- Minimalist CLI environnement, Arch Linux based on Linux 4.14.13-1.
- Three compiler analysed: GCC, ICC and CLang (all up to date).
- Input file generated by playing the program against itself, deterministic program.
- The executions are repeated 20 times for each tested configuration.

Analysis and optimization of sequential performances



Code performance profiling

GProf

% time	cumulative seconds	calls	total s/call	name
48.83	167.72	5172852992	0.00	jouer_coup(Next*, Pos*, Pos*, int, int)
18.25	230.39	3262157329	0.00	copier(Pos*, Pos*)
17.22	289.56	3262157262	0.00	est_affame(Pos*, int)
8.79	319.75	1457826958	0.00	calculer_coup(Next*, Pos*, int, int, int, int, bool)
5.31	337.98	3252245939	0.00	valeur_minimaxAB(Next*, Pos*, int, int, int, int, bool)
0.87	340.99	1775297034	0.00	evaluer(Pos*)
0.52	342.77	67	0.03	test_fin(Pos*)
0.09	343.65	33	10.34	decisionAB(Next*, Pos*, int, bool)

Figure: Code Profiling with GProf.

Code performance profiling

vTune

Elapsed Time [?]: 118.666s

➤ CPU Time [?] :	117.755s
Instructions Retired:	588,377,500,000
CPI Rate [?] :	0.499
CPU Frequency Ratio [?] :	1.000
Total Thread Count:	1
Paused Time [?] :	0s

Top Hotspots

This section lists the most active functions in your application.

Function	Module	CPU Time [?]
jouer_coup	advisor.out	70.132s
calculer_coup	advisor.out	15.441s
copier	advisor.out	14.165s
est_affame	advisor.out	7.176s
valeur_minimaxAB	advisor.out	5.755s
[Others]		5.087s


Elapsed Time [?]: 120.174s

CPU Time [?] :	118.454s
▼ Memory Bound [?] :	2.7%
L1 Bound [?] :	6.6%
▼ DRAM Bound [?] :	
DRAM Bandwidth Bound [?] :	0.0%
LLC Miss [?] :	0.1%
Loads:	157,360,120,662
Stores:	100,867,825,944
LLC Miss Count [?] :	900,054
Average Latency (cycles) [?] :	8
Total Thread Count:	1
Paused Time [?] :	0s

Figure: vTune memory usage analysis.

Figure: vTune processor usage analysis.

Naïve approach to adding parallelism



False Sharing

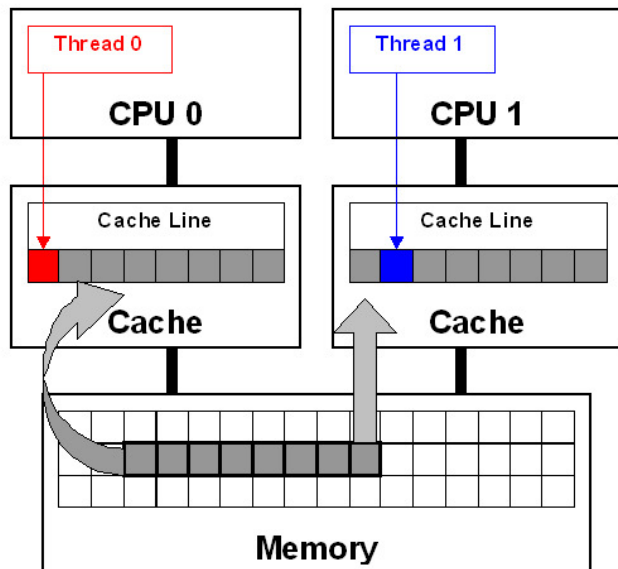



Figure: False Sharing.

Source : Intel's Documentation.

Thread Affinity

Figure: Analyzing the impact of thread placement with the Intel compiler, **-O3** with 8 CPUs.



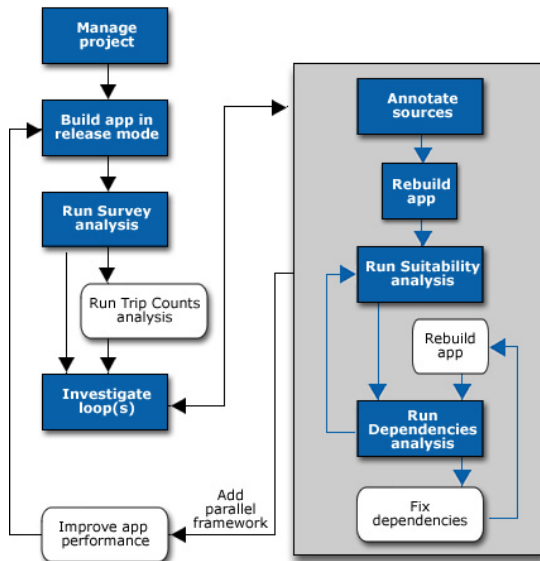


Figure: Iterative method to add parallelism (proposed by Intel).

Source : Intel's Documentation.

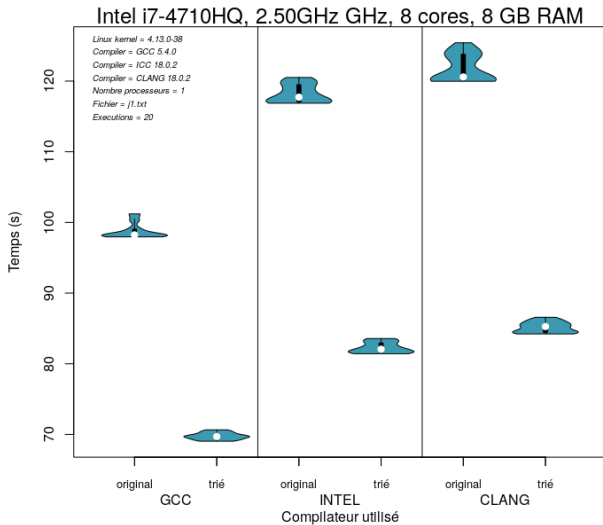
Difficulties of code optimization

- Very sequential algorithm.
- Low number of iterations ($n = 6$).
- Lot of usage of pointers that could blind the compiler's optimisation passes.

Simple algorithmic modification

- Sort the "moves" to evaluate.
- Transposition table.

Sort the "moves"



Transposition table

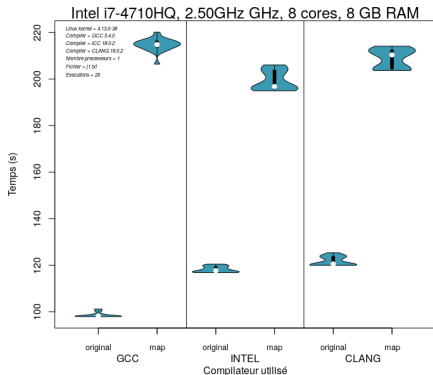


Figure: Binary search tree.

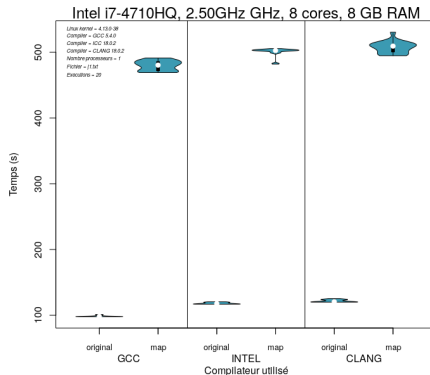


Figure: Hash table.

- GCC beats ICC for a sequential code on Intel's architecture.
- Parallelization with OpenMP can severely degrade performance.
- False Sharing phenomena.
- Algorithmic changes needed to improve performance.

- Array padding to reduce the false sharing problem.
- Speculative execution.